

Peer Review Final Assignment

Introduction

In this lab, you will build an image classifier using the VGG16 pre-trained model, and you will evaluate it and compare its performance to the model we built in the last module using the ResNet50 pre-trained model. Good luck!

Download Data

Use the `wget` command to download the data for this assignment from here: https://s3-api.us-gio.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/DL0321EN/data/concrete_data_week4.zip (https://s3-api.us-gio.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/DL0321EN/data/concrete_data_week4.zip)

Use the following cells to download the data.

```
In [ ]: #!/wget https://s3-api.us-gio.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/DL0321EN/data/concrete_data_week4.zip
```

```
In [ ]: #!/zip concrete_data_week4.zip
```

After you unzip the data, you will find the data has already been divided into a train, validation, and test sets.

Part 1

In this part, you will design a classifier using the VGG16 pre-trained model. Just like the ResNet50 model, you can import the model `VGG16` from `keras.applications`.

You will essentially build your classifier as follows:

1. Import libraries, modules, and packages you will need. Make sure to import the `preprocess_input` function from `keras.applications.vgg16`.
2. Use a batch size of 100 images for both training and validation.
3. Construct an `ImageDataGenerator` for the training set and another one for the validation set. VGG16 was originally trained on 224×224 images, so make sure to address that when defining the `ImageDataGenerator` instances.
4. Create a sequential model using Keras. Add VGG16 model to it and dense layer.
5. Compile the model using the adam optimizer and the categorical_crossentropy loss function.
6. Fit the model on the augmented data using the `ImageDataGenerators`.

Use the following cells to create your classifier.

```
In [1]: import numpy as np
import pandas as pd
%matplotlib inline
import matplotlib
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras

from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.applications import VGG16
from tensorflow.keras.applications.vgg16 import preprocess_input
from tensorflow.keras import models
from tensorflow.keras import layers
from tensorflow.keras import optimizers
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
from tensorflow.keras.models import load_model
```

```
In [2]: train_dir = r'C:\Users\Dennis\Desktop\AI Capstone Project with Deep Learning\con
crete_data_week4\train'
validation_dir = r'C:\Users\Dennis\Desktop\AI Capstone Project with Deep Learnin
g\concrete_data_week4\valid'
test_dir = r'C:\Users\Dennis\Desktop\AI Capstone Project with Deep Learning\concr
ete_data_week4\test'
```

```
In [3]: # Generating batches of tensor image data
train_datagen = ImageDataGenerator(rescale=1./255)
valid_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(224, 224),
    batch_size=100,
    class_mode='categorical')

validation_generator = valid_datagen.flow_from_directory(
    validation_dir,
    target_size=(224, 224),
    batch_size=100,
    class_mode='categorical')
```

Found 30000 images belonging to 2 classes.
Found 9500 images belonging to 2 classes.

```
In [4]: conv_base = VGG16(weights='imagenet',
    include_top=False,
    input_shape=(224, 224, 3))
```

```
In [5]: conv_base.summary()
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
<hr/>		
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
<hr/>		
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
<hr/>		
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
<hr/>		
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
<hr/>		
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
<hr/>		
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
<hr/>		
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
<hr/>		
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
<hr/>		
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
<hr/>		
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
<hr/>		
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
<hr/>		
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
<hr/>		
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
<hr/>		
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
<hr/>		
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
<hr/>		
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
<hr/>		
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
<hr/>		
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

```
In [6]: conv_base.trainable = False
```

```
In [7]: model = models.Sequential()  
model.add(conv_base)  
model.add(layers.Flatten())  
model.add(layers.Dense(2,activation='softmax'))
```

```
In [8]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====	=====	=====
vgg16 (Model)	(None, 7, 7, 512)	14714688
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 2)	50178
=====	=====	=====

Total params: 14,764,866
Trainable params: 50,178
Non-trainable params: 14,714,688

```
In [9]: model.compile(optimizer='adam',loss='categorical_crossentropy',metrics=['accuracy'])
```

```
In [10]: checkpoint = ModelCheckpoint("vgg16.h5",monitor='val_loss',verbose=1,save_best_only=True)
```

```
In [11]: early = EarlyStopping(monitor='val_loss',min_delta=0,patience=3,verbose=1)
```

```
In [12]: history = model.fit_generator(generator=train_generator, steps_per_epoch=3, epochs=5, verbose=1,
                                       validation_data=validation_generator, validation_steps=5, callbacks=[checkpoint, early])
```

Epoch 1/5

2/3 [=====>.....] - ETA: 2:51 - loss: 0.9101 - accuracy: 0.5800

Epoch 00001: val_loss improved from inf to 0.51783, saving model to vgg16.h5
WARNING:tensorflow:Early stopping conditioned on metric `val_acc` which is not available. Available metrics are: loss, accuracy, val_loss, val_accuracy

3/3 [=====] - 767s 256s/step - loss: 0.7215 - accuracy: 0.6967 - val_loss: 0.5178 - val_accuracy: 0.5780

Epoch 2/5

2/3 [=====>.....] - ETA: 2:51 - loss: 0.4982 - accuracy: 0.6800

Epoch 00002: val_loss improved from 0.51783 to 0.19915, saving model to vgg16.h5

WARNING:tensorflow:Early stopping conditioned on metric `val_acc` which is not available. Available metrics are: loss, accuracy, val_loss, val_accuracy

3/3 [=====] - 783s 261s/step - loss: 0.4007 - accuracy: 0.7733 - val_loss: 0.1991 - val_accuracy: 0.9420

Epoch 3/5

2/3 [=====>.....] - ETA: 2:57 - loss: 0.2160 - accuracy: 0.9300

Epoch 00003: val_loss did not improve from 0.19915

WARNING:tensorflow:Early stopping conditioned on metric `val_acc` which is not available. Available metrics are: loss, accuracy, val_loss, val_accuracy

3/3 [=====] - 755s 252s/step - loss: 0.2165 - accuracy: 0.9167 - val_loss: 0.2091 - val_accuracy: 0.9080

Epoch 4/5

2/3 [=====>.....] - ETA: 2:35 - loss: 0.1507 - accuracy: 0.9400

Epoch 00004: val_loss improved from 0.19915 to 0.11376, saving model to vgg16.h5

WARNING:tensorflow:Early stopping conditioned on metric `val_acc` which is not available. Available metrics are: loss, accuracy, val_loss, val_accuracy

3/3 [=====] - 704s 235s/step - loss: 0.1282 - accuracy: 0.9567 - val_loss: 0.1138 - val_accuracy: 0.9620

Epoch 5/5


2/3 [=====>.....] - ETA: 2:25 - loss: 0.0886 - accuracy: 0.9650

Epoch 00005: val_loss did not improve from 0.11376

WARNING:tensorflow:Early stopping conditioned on metric `val_acc` which is not available. Available metrics are: loss, accuracy, val_loss, val_accuracy

3/3 [=====] - 709s 236s/step - loss: 0.1406 - accuracy: 0.9500 - val_loss: 0.1304 - val_accuracy: 0.9500

```
In [13]: #Plot Graph to see the result
plt.figure(figsize=(8,5))
plt.plot(history.history["accuracy"])
plt.plot(history.history['val_accuracy'])
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title("Model Accuracy")
plt.ylabel("Accuracy")
plt.xlabel("Epoch")
plt.legend(["Accuracy", "Validation Accuracy", "Loss", "Validation Loss"])
plt.show()
```



```
In [14]: #Save the model
model.save('vgg16.h5')
```

```
In [16]: del model
```

```
In [17]: vgg16 = load_model('vgg16.h5')
```

```
In [18]: vgg16.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
vgg16 (Model)	(None, 7, 7, 512)	14714688
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 2)	50178
Total params: 14,764,866		
Trainable params: 50,178		
Non-trainable params: 14,714,688		

Part 2

In this part, you will evaluate your deep learning models on a test data. For this part, you will need to do the following:

1. Load your saved model that was built using the ResNet50 model.
2. Construct an ImageDataGenerator for the test set. For this ImageDataGenerator instance, you only need to pass the directory of the test images, target size, and the **shuffle** parameter and set it to False.
3. Use the **evaluate_generator** method to evaluate your models on the test data, by passing the above ImageDataGenerator as an argument. You can learn more about **evaluate_generator** [here \(https://keras.io/models/sequential/\)](https://keras.io/models/sequential/).
4. Print the performance of the classifier using the VGG16 pre-trained model.
5. Print the performance of the classifier using the ResNet pre-trained model.

Use the following cells to evaluate your models.

```
In [19]: resnet = load_model("resnet.h5")
```

WARNING:tensorflow:Error in loading the saved optimizer state. As a result, your model is starting with a freshly initialized optimizer.

```
In [20]: resnet.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====	=====	=====
resnet50 (Model)	(None, 2048)	23587712
dense_1 (Dense)	(None, 2)	4098
=====	=====	=====
Total params: 23,591,810		
Trainable params: 23,538,690		
Non-trainable params: 53,120		
=====		

```
In [21]: test_datagen = ImageDataGenerator(rescale=1./255)

test_generator = test_datagen.flow_from_directory(
    test_dir,
    target_size=(224, 224),
    batch_size=100,
    class_mode='categorical',
    shuffle=False)
```

Found 500 images belonging to 2 classes.

```
In [22]: len(test_generator)
```

```
Out[22]: 5
```

```
In [23]: loss = resnet.evaluate_generator(test_generator, steps=5, verbose=1)
```

5/5 [=====] - 838s 168s/step - loss: 1.3265 - accuracy: 0.5000

```
In [24]: loss
```

```
Out[24]: [1.3265284791588783, 0.5]
```

```
In [25]: resnet.metrics_names
```

```
Out[25]: ['loss', 'accuracy']
```

```
In [27]: print("Accuracy for ResNet is %.2f%%" % (loss[1]*100))
```

```
Accuracy for ResNet is 50.00%
```

```
In [28]: loss2 = vgg16.evaluate_generator(test_generator, steps=5, verbose=1)
```

```
5/5 [=====] - 367s 73s/step - loss: 0.1464 - accuracy: 0.9480
```

```
In [29]: print("Accuracy for VGG is %.2f%%" % (loss2[1]*100))
```

```
Accuracy for VGG is 94.80%
```

Part 3

In this model, you will predict whether the images in the test data are images of cracked concrete or not. You will do the following:

1. Use the **predict_generator** method to predict the class of the images in the test data, by passing the test data ImageDataGenerator instance defined in the previous part as an argument. You can learn more about the **predict_generator** method [here \(https://keras.io/models/sequential/\)](https://keras.io/models/sequential/).
2. Report the class predictions of the first five images in the test set. You should print something list this:

```
Positive  
Negative  
Positive  
Positive  
Negative
```

Use the following cells to make your predictions.

```
In [30]: resnet_predict = resnet.predict_generator(generator=test_generator, steps=5, verbose=1)
```

```
5/5 [=====] - 844s 169s/step
```



```
In [31]: resnet_predict
```

```
Out[31]: array([[0.9216869 , 0.07831309],
 [0.92157716, 0.07842289],
 [0.9215356 , 0.07846433],
 [0.92253685, 0.07746314],
 [0.9214253 , 0.07857466],
 [0.92207825, 0.0779217 ],
 [0.9212901 , 0.07870995],
 [0.92132145, 0.07867851],
 [0.9213929 , 0.0786071 ],
 [0.92163324, 0.07836676],
 [0.9213261 , 0.07867392],
 [0.9213351 , 0.07866489],
 [0.92126745, 0.07873252],
 [0.9213651 , 0.07863495],
 [0.92131376, 0.07868626],
 [0.921574 , 0.07842605],
 [0.9212273 , 0.07877272],
 [0.9214263 , 0.07857364],
 [0.92137116, 0.07862884],
 [0.92178464, 0.07821533],
 [0.9210988 , 0.07890116],
 [0.92114717, 0.07885284],
 [0.9213146 , 0.07868544],
 [0.92123556, 0.07876439],
 [0.92121005, 0.07878993],
 [0.92123306, 0.07876692],
 [0.9215979 , 0.07840209],
 [0.92157644, 0.0784236 ],
 [0.9213207 , 0.07867935],
 [0.92119765, 0.0788024 ],
 [0.92277926, 0.07722078],
 [0.92114717, 0.07885284],
 [0.92187744, 0.07812262],
 [0.9210526 , 0.07894742],
 [0.9212905 , 0.07870954],
 [0.9209829 , 0.0790171 ],
 [0.9214592 , 0.07854079],
 [0.9231534 , 0.07684664],
 [0.9210671 , 0.07893284],
 [0.92161804, 0.0783819 ],
 [0.9214698 , 0.07853014],
 [0.9212276 , 0.07877243],
 [0.92099243, 0.07900761],
 [0.9221099 , 0.07789013],
 [0.9216296 , 0.07837044],
 [0.9214747 , 0.07852531],
 [0.92124903, 0.07875095],
 [0.9221125 , 0.07788745],
 [0.9228376 , 0.07716238],
 [0.92120045, 0.07879959],
 [0.92253053, 0.07746945],
 [0.92123544, 0.07876457],
 [0.92119914, 0.07880089],
 [0.9212064 , 0.07879357],
 [0.9210919 , 0.07890806],
 [0.9210202 , 0.07897975],
 [0.92116696, 0.07883306],
 [0.92123294, 0.07876705],
 [0.92111355, 0.07888643],
 [0.92140496, 0.078595 ],
 [0.92322266, 0.07677736],
 [0.9213832 , 0.07861682],
 [0.9211256 , 0.07887437],
 [0.9211652 , 0.07883476],
 [0.9212341 , 0.07876594],
 [0.9217736 , 0.07822638],
 [0.92124444, 0.07875548],
```

```
In [32]: len(resnet_predict)
```

```
Out[32]: 500
```

```
In [33]: np.round(a=resnet_predict,decimals=3)
```

```
Out[33]: array([[0.922, 0.078],
 [0.922, 0.078],
 [0.922, 0.078],
 [0.923, 0.077],
 [0.921, 0.079],
 [0.922, 0.078],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.922, 0.078],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.922, 0.078],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.922, 0.078],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.922, 0.078],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.923, 0.077],
 [0.921, 0.079],
 [0.922, 0.078],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.923, 0.077],
 [0.921, 0.079],
 [0.922, 0.078],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.922, 0.078],
 [0.922, 0.078],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.922, 0.078],
 [0.923, 0.077],
 [0.921, 0.079],
 [0.923, 0.077],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.923, 0.077],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.921, 0.079],
 [0.922, 0.078],
 [0.921, 0.079]
```

```
In [34]: vgg_predict = vgg16.predict_generator(generator=test_generator, steps=5, verbose=1)
```

```
5/5 [=====] - 327s 65s/step
```

```
In [35]: vgg_predict
```

```
Out[35]: array([[9.78436530e-01, 2.15634406e-02],
 [6.47004917e-02, 9.35299516e-01],
 [9.66751296e-03, 9.90332425e-01],
 [2.45390132e-01, 7.54609883e-01],
 [9.77425218e-01, 2.25748252e-02],
 [9.20497298e-01, 7.95026794e-02],
 [9.48500633e-01, 5.14993444e-02],
 [9.75193679e-01, 2.48063263e-02],
 [9.78029191e-01, 2.19708458e-02],
 [9.76324797e-01, 2.36752164e-02],
 [9.71785069e-01, 2.82149743e-02],
 [9.45410550e-01, 5.45894057e-02],
 [9.78415310e-01, 2.15847194e-02],
 [9.77163970e-01, 2.28359848e-02],
 [6.42237902e-01, 3.57762069e-01],
 [9.62787509e-01, 3.72124314e-02],
 [9.71709430e-01, 2.82905605e-02],
 [9.80994999e-01, 1.90049447e-02],
 [9.67856646e-01, 3.21433507e-02],
 [1.85327381e-01, 8.14672709e-01],
 [9.58345711e-01, 4.16543335e-02],
 [9.65750992e-01, 3.42490077e-02],
 [9.80306804e-01, 1.96931977e-02],
 [9.79576111e-01, 2.04238780e-02],
 [9.56659794e-01, 4.33402695e-02],
 [9.57621813e-01, 4.23782282e-02],
 [5.33222079e-01, 4.66777891e-01],
 [9.30389881e-01, 6.96101189e-02],
 [6.24859929e-01, 3.75140041e-01],
 [9.60869491e-01, 3.91304828e-02],
 [2.80978769e-01, 7.19021261e-01],
 [9.50272262e-01, 4.97277938e-02],
 [6.23670995e-01, 3.76329035e-01],
 [9.67287123e-01, 3.27128880e-02],
 [9.71206188e-01, 2.87937596e-02],
 [9.55041289e-01, 4.49587218e-02],
 [2.69104898e-01, 7.30895102e-01],
 [2.06871420e-01, 7.93128550e-01],
 [9.79833364e-01, 2.01667026e-02],
 [9.52963948e-01, 4.70360070e-02],
 [9.77202952e-01, 2.27971021e-02],
 [7.35875249e-01, 2.64124751e-01],
 [9.83925223e-01, 1.60747562e-02],
 [9.22190309e-01, 7.78096765e-02],
 [3.51966619e-01, 6.48033381e-01],
 [9.68531668e-01, 3.14682610e-02],
 [9.60019171e-01, 3.99808511e-02],
 [3.14741373e-01, 6.85258567e-01],
 [1.17344216e-01, 8.82655799e-01],
 [9.68562245e-01, 3.14378217e-02],
 [1.54721618e-01, 8.45278382e-01],
 [9.73627508e-01, 2.63725519e-02],
 [9.73936200e-01, 2.60638446e-02],
 [9.78680611e-01, 2.13193540e-02],
 [9.73119557e-01, 2.68804152e-02],
 [9.55021381e-01, 4.49786484e-02],
 [9.73639727e-01, 2.63603479e-02],
 [9.81054425e-01, 1.89456120e-02],
 [8.04405630e-01, 1.95594415e-01],
 [9.60171878e-01, 3.98281068e-02],
 [3.47396642e-01, 6.52603388e-01],
 [9.74164665e-01, 2.58352850e-02],
 [9.79163945e-01, 2.08360404e-02],
 [9.54891562e-01, 4.51083779e-02],
 [9.76826370e-01, 2.31735520e-02],
 [7.46428847e-01, 2.53571153e-01],
 [9.79618669e-01, 2.03813501e-02],
```



```
In [36]: len(vgg_predict)
```

```
Out[36]: 500
```

```
In [37]: np.round(a=vgg_predict,decimals=3)
```

```
Out[37]: array([[0.978, 0.022],
                [0.065, 0.935],
                [0.01 , 0.99 ],
                [0.245, 0.755],
                [0.977, 0.023],
                [0.92 , 0.08 ],
                [0.949, 0.051],
                [0.975, 0.025],
                [0.978, 0.022],
                [0.976, 0.024],
                [0.972, 0.028],
                [0.945, 0.055],
                [0.978, 0.022],
                [0.977, 0.023],
                [0.642, 0.358],
                [0.963, 0.037],
                [0.972, 0.028],
                [0.981, 0.019],
                [0.968, 0.032],
                [0.185, 0.815],
                [0.958, 0.042],
                [0.966, 0.034],
                [0.98 , 0.02 ],
                [0.98 , 0.02 ],
                [0.957, 0.043],
                [0.958, 0.042],
                [0.533, 0.467],
                [0.93 , 0.07 ],
                [0.625, 0.375],
                [0.961, 0.039],
                [0.281, 0.719],
                [0.95 , 0.05 ],
                [0.624, 0.376],
                [0.967, 0.033],
                [0.971, 0.029],
                [0.955, 0.045],
                [0.269, 0.731],
                [0.207, 0.793],
                [0.98 , 0.02 ],
                [0.953, 0.047],
                [0.977, 0.023],
                [0.736, 0.264],
                [0.984, 0.016],
                [0.922, 0.078],
                [0.352, 0.648],
                [0.969, 0.031],
                [0.96 , 0.04 ],
                [0.315, 0.685],
                [0.117, 0.883],
                [0.969, 0.031],
                [0.155, 0.845],
                [0.974, 0.026],
                [0.974, 0.026],
                [0.979, 0.021],
                [0.973, 0.027],
                [0.955, 0.045],
                [0.974, 0.026],
                [0.981, 0.019],
                [0.804, 0.196],
                [0.96 , 0.04 ],
                [0.347, 0.653],
                [0.974, 0.026],
                [0.979, 0.021],
                [0.955, 0.045],
                [0.977, 0.023],
                [0.746, 0.254],
                [0.98 , 0.02 ],
```

```
In [38]: classes = np.round(a=vgg_predict,decimals=3)
```

```
In [39]: classes
```

```
Out[39]: array([[0.978, 0.022],
                [0.065, 0.935],
                [0.01 , 0.99 ],
                [0.245, 0.755],
                [0.977, 0.023],
                [0.92 , 0.08 ],
                [0.949, 0.051],
                [0.975, 0.025],
                [0.978, 0.022],
                [0.976, 0.024],
                [0.972, 0.028],
                [0.945, 0.055],
                [0.978, 0.022],
                [0.977, 0.023],
                [0.642, 0.358],
                [0.963, 0.037],
                [0.972, 0.028],
                [0.981, 0.019],
                [0.968, 0.032],
                [0.185, 0.815],
                [0.958, 0.042],
                [0.966, 0.034],
                [0.98 , 0.02 ],
                [0.98 , 0.02 ],
                [0.957, 0.043],
                [0.958, 0.042],
                [0.533, 0.467],
                [0.93 , 0.07 ],
                [0.625, 0.375],
                [0.961, 0.039],
                [0.281, 0.719],
                [0.95 , 0.05 ],
                [0.624, 0.376],
                [0.967, 0.033],
                [0.971, 0.029],
                [0.955, 0.045],
                [0.269, 0.731],
                [0.207, 0.793],
                [0.98 , 0.02 ],
                [0.953, 0.047],
                [0.977, 0.023],
                [0.736, 0.264],
                [0.984, 0.016],
                [0.922, 0.078],
                [0.352, 0.648],
                [0.969, 0.031],
                [0.96 , 0.04 ],
                [0.315, 0.685],
                [0.117, 0.883],
                [0.969, 0.031],
                [0.155, 0.845],
                [0.974, 0.026],
                [0.974, 0.026],
                [0.979, 0.021],
                [0.973, 0.027],
                [0.955, 0.045],
                [0.974, 0.026],
                [0.981, 0.019],
                [0.804, 0.196],
                [0.96 , 0.04 ],
                [0.347, 0.653],
                [0.974, 0.026],
                [0.979, 0.021],
                [0.955, 0.045],
                [0.977, 0.023],
                [0.746, 0.254],
                [0.98 , 0.02 ],
```

```
In [40]: filenames=test_generator.filenames  
filenames
```

```
Out[40]: ['negative\\19751.jpg',
'negative\\19752.jpg',
'negative\\19753.jpg',
'negative\\19754.jpg',
'negative\\19755.jpg',
'negative\\19756.jpg',
'negative\\19757.jpg',
'negative\\19758.jpg',
'negative\\19759.jpg',
'negative\\19760.jpg',
'negative\\19761.jpg',
'negative\\19762.jpg',
'negative\\19763.jpg',
'negative\\19764.jpg',
'negative\\19765.jpg',
'negative\\19766.jpg',
'negative\\19767.jpg',
'negative\\19768.jpg',
'negative\\19769.jpg',
'negative\\19770.jpg',
'negative\\19771.jpg',
'negative\\19772.jpg',
'negative\\19773.jpg',
'negative\\19774.jpg',
'negative\\19775.jpg',
'negative\\19776.jpg',
'negative\\19777.jpg',
'negative\\19778.jpg',
'negative\\19779.jpg',
'negative\\19780.jpg',
'negative\\19781.jpg',
'negative\\19782.jpg',
'negative\\19783.jpg',
'negative\\19784.jpg',
'negative\\19785.jpg',
'negative\\19786.jpg',
'negative\\19787.jpg',
'negative\\19788.jpg',
'negative\\19789.jpg',
'negative\\19790.jpg',
'negative\\19791.jpg',
'negative\\19792.jpg',
'negative\\19793.jpg',
'negative\\19794.jpg',
'negative\\19795.jpg',
'negative\\19796.jpg',
'negative\\19797.jpg',
'negative\\19798.jpg',
'negative\\19799.jpg',
'negative\\19800.jpg',
'negative\\19801.jpg',
'negative\\19802.jpg',
'negative\\19803.jpg',
'negative\\19804.jpg',
'negative\\19805.jpg',
'negative\\19806.jpg',
'negative\\19807.jpg',
'negative\\19808.jpg',
'negative\\19809.jpg',
'negative\\19810.jpg',
'negative\\19811.jpg',
'negative\\19812.jpg',
'negative\\19813.jpg',
'negative\\19814.jpg',
'negative\\19815.jpg',
'negative\\19816.jpg',
'negative\\19817.jpg',
```



```
In [41]: results = pd.DataFrame({"file": filenames, "prediction": vgg_predict[:, 0], "class":  
classes[:, 0]})
```

```
In [42]: results
```

Out [42]:

	file	prediction	class
0	negative\19751.jpg	0.978437	0.978
1	negative\19752.jpg	0.064700	0.065
2	negative\19753.jpg	0.009668	0.010
3	negative\19754.jpg	0.245390	0.245
4	negative\19755.jpg	0.977425	0.977
...
495	positive\19996.jpg	0.001427	0.001
496	positive\19997.jpg	0.190317	0.190
497	positive\19998.jpg	0.005833	0.006
498	positive\19999.jpg	0.001227	0.001
499	positive\20000.jpg	0.042057	0.042

500 rows × 3 columns

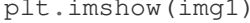
Assumption is prediction > 0.5 , class is 1 else prediction < 0.5, class is 0

Method 2: Sampling 5 images randomly and predict class

```
In [48]: from tensorflow.keras.preprocessing import image
```

```
In [62]: img1 = image.load_img("19751.jpg",target_size=(224,224))

# img1 = tf.cast(img1, tf.float32)
img1 = np.asarray(img1)

plt.imshow(img1)

#img1 = np.expand_dims(img1, axis=0)

output1 = vgg16.predict(img1)
print(output1)

if output1[0][1] == 0:
    print("negative")
else:
    print('positive')
```

```

-----
ValueError                                Traceback (most recent call last)
<ipython-input-62-ec2468f651f1> in <module>
      7 #img1 = np.expand_dims(img1, axis=0)
      8
----> 9 output1 = vgg16.predict(img1)
     10 print(output1)
     11

C:\ProgramData\Anaconda3\lib\site-packages\tensorflow_core\python\keras\engin
e\training.py in predict(self, x, batch_size, verbose, steps, callbacks, max_q
ueue_size, workers, use_multiprocessing)
     907         max_queue_size=max_queue_size,
     908         workers=workers,
--> 909         use_multiprocessing=use_multiprocessing)
     910
     911     def reset_metrics(self):

C:\ProgramData\Anaconda3\lib\site-packages\tensorflow_core\python\keras\engin
e\training_v2.py in predict(self, model, x, batch_size, verbose, steps, callba
cks, **kwargs)
     460         return self._model_iteration(
     461             model, ModeKeys.PREDICT, x=x, batch_size=batch_size, verbose=v
erbose,
--> 462             steps=steps, callbacks=callbacks, **kwargs)
     463
     464


C:\ProgramData\Anaconda3\lib\site-packages\tensorflow_core\python\keras\engin
e\training_v2.py in _model_iteration(self, model, mode, x, y, batch_size, verb
ose, sample_weight, steps, callbacks, **kwargs)
     394         sample_weights=sample_weight,
     395         steps=steps,
--> 396         distribution_strategy=strategy)
     397         total_samples = _get_total_number_of_samples(adapter)
     398         use_sample = total_samples is not None

C:\ProgramData\Anaconda3\lib\site-packages\tensorflow_core\python\keras\engin
e\training_v2.py in _process_inputs(model, x, y, batch_size, epochs, sample_we
ights, class_weights, shuffle, steps, distribution_strategy, max_queue_size, w
orkers, use_multiprocessing)
     592         batch_size=batch_size,
     593         check_steps=False,
--> 594         steps=steps)
     595         adapter = adapter_cls(
     596             x,

C:\ProgramData\Anaconda3\lib\site-packages\tensorflow_core\python\keras\engin
e\training.py in _standardize_user_data(self, x, y, sample_weight, class_weigh
t, batch_size, check_steps, steps_name, steps, validation_split, shuffle, extr
act_tensors_from_dataset)
     2470         feed_input_shapes,
     2471         check_batch_axis=False, # Don't enforce the batch size.
-> 2472         exception_prefix='input')
     2473
     2474         # Get typespecs for the input data and sanitize it if necessary.

C:\ProgramData\Anaconda3\lib\site-packages\tensorflow_core\python\keras\engin
e\training_utils.py in standardize_input_data(data, names, shapes, check_batch
_axis, exception_prefix)
     563         ': expected ' + names[i] + ' to have ' +
     564         str(len(shape)) + ' dimensions, but got arr
ay '
--> 565         'with shape ' + str(data_shape))
     566         if not check_batch_axis:
     567             data_shape = data_shape[1:]

```



In []:

In []:

In []:

In []:

Thank you for completing this lab!

This notebook was created by Alex Aklson.

This notebook is part of a course on **Coursera** called *AI Capstone Project with Deep Learning*. If you accessed this notebook outside the course, you can take this course online by clicking [here \(https://cocl.us/DL0321EN_Coursera_Week4_LAB1\)](https://cocl.us/DL0321EN_Coursera_Week4_LAB1).

Copyright © 2020 [IBM Developer Skills Network \(https://cognitiveclass.ai/?utm_source=bducopyrightlink&utm_medium=dswb&utm_campaign=bdu\)](https://cognitiveclass.ai/?utm_source=bducopyrightlink&utm_medium=dswb&utm_campaign=bdu). This notebook and its source code are released under the terms of the [MIT License \(https://bigdatauniversity.com/mit-license/\)](https://bigdatauniversity.com/mit-license/).